News Impact Modeling Using a Two-Stage XGBoost Architecture for Stock Forecasting

Dushan Warnakulasuriya

(IT22069672)

BSc (Hons) degree in Information Technology Specializing in Software Engineering

Department of Information Technology

Sri Lanka Institute of Information Technology
Sri Lanka

August 2026

News Impact Modeling Using a Two-Stage XGBoost Architecture for Stock Forecasting

Dushan Warnakulasuriya

(IT22069672)

BSc (Hons) degree in Information Technology Specializing in Software Engineering

Department of Information Technology

Sri Lanka Institute of Information Technology
Sri Lanka

August 2026

DECLARATION OF THE CANDIDATE AND SUPERVISOR

I declare that this is my own work and this proposal does not incorporate without acknowledgment any material previously submitted for a degree or diploma in any other university or Institute of higher learning and to the best of my knowledge and belief it does not contain any material previously published or written by another person except where the acknowledgment is made in the text.

Name	Student ID	Signature
Bolonghe W D U	IT22069672	

The supervisor/s should certify the proposal report with the following declaration.

The above candidate is carrying out research for the undergraduate Dissertation under my supervision.

…….. ……
Signature of the supervisor: Date

…….. ……
Signature of the co-supervisor: Date

Abstract

This study presents a sentiment analysis framework designed to extract and quantify market-relevant information from financial textual data. In modern financial ecosystems, unstructured data sources such as news articles, corporate disclosures, and financial reports significantly influence investor behavior and market movements. The objective of this research is to convert such qualitative information into structured sentiment indicators that can provide actionable insights. The proposed methodology involves systematic data collection followed by preprocessing steps, including tokenization, stop-word removal, and text normalization. Natural Language Processing (NLP) techniques are applied to transform textual content into numerical representations. A hybrid sentiment classification approach is implemented, combining lexicon-based methods with supervised machine learning algorithms to categorize text into positive, negative, or neutral sentiment classes. The model is trained and validated using labeled datasets to ensure reliability and generalization. Experimental results demonstrate that the developed framework effectively captures underlying sentiment patterns within financial text and identifies shifts in market perception. The findings indicate that sentiment analysis can serve as a reliable tool for interpreting qualitative financial information and enhancing decision-making processes. In conclusion, this research highlights the significance of sentiment analysis in financial text mining and its potential to generate meaningful market insights. Future work should explore the integration of advanced deep learning techniques and real-time data streams to further improve accuracy and responsiveness.

Keywords: ***Sentiment analysis, financial text mining, natural language processing, machine learning, market sentiment***

Acknowledgement

TABLE OF CONTENTS

DECLARATION OF THE CANDIDATE & SUPERVISOR iii

ABSTRACT iv

ACKNOWLEDGEMENT v

TABLE OF CONTENTS vi

LIST OF FIGURES viii

LIST OF TABLES xi

LIST OF ABBREVIATIONS xii

1. INTRODUCTION 1

1.1. Background and Literature survey 4

1.2. Research Gap 10

1.3. Research Problem 12

1.4. Research Objectives 14

1.4.1 Main Objective 14

1.4.2 Specific Objectives 14

2. METHODOLOGY 16

2.1. Understanding Key pillars of the research domain 16

2.1.1 Deep Learning 16

2.1.2 Convolutional Neural Network (CNN) 16

2.1.3 Transfer Learning based on deep CNN 17

2.1.4 Data Augmentation 18

2.2. Approach 20

2.2.1 Data collection 21

2.2.2 Methodology 23

2.2.3 Transfer learning vs. Non-Transfer learning 24

2.3. Summary of the methodology 28

2.4. High level System Architecture Diagram 29

2.5. Project Requirements that have been achieved 30

2.5.1 Functional Requirements 30
2.5.2 Non-Functional Requirements 31
2.5.3 Other Requirements 31
2.6. Consideration of aspects of the system 32
2.6.1 Social aspects 32
2.6.2 Security aspects 32
2.6.3 Ethical aspects 33
2.6.4 Limitations 33
2.7. Commercialization aspects of the product 34
2.7.1 Target Audience 34
2.7.2 Market Space 34
2.7.3 Revenue Earning 34
2.8. Testing and Implementation 35
2.8.1 Code Implementation of the research part 35
2.8.2 Testing 51
2.9. Tools and Technologies 53
3. RESULTS AND DISCUSSION 54
3.1. Results 54
3.2. Research Findings and Discussion 61
4. CONCLUSION 63
REFERENCES 64

List of Figures

List of Tables

List of Abbreviations

List of Appendices

1. INTRODUCTION

Financial markets operate as complex, information-driven systems in which prices are continuously influenced by the flow and interpretation of new information. Traditionally, financial analysis has focused on structured quantitative data such as historical prices, financial ratios, earnings reports, and macroeconomic indicators. While these factors remain fundamental, they do not fully capture the behavioral and psychological elements that significantly affect market dynamics. A growing body of evidence suggests that investor sentiment—shaped by opinions, expectations, and reactions to news—plays a critical role in explaining short-term market movements and volatility. Much of this sentiment is embedded in unstructured textual data, including financial news articles, corporate announcements, earnings call transcripts, and social media discussions.

With the rapid expansion of digital information sources, the volume of financial text data has increased exponentially, making manual analysis impractical and inefficient. As a result, automated techniques for extracting meaningful insights from textual data have become essential. Sentiment analysis, a specialized area within Natural Language Processing (NLP), offers a systematic approach to identifying and quantifying subjective information in text. By classifying textual content into sentiment categories such as positive, negative, or neutral, sentiment analysis enables the transformation of qualitative narratives into structured, machine-readable features.

In the context of financial markets, sentiment analysis serves as a bridge between qualitative information and quantitative modeling. It allows analysts to capture market perception, investor confidence, and reaction to events, which are often not directly observable through numerical data alone. For instance, a company's financial results may meet expectations, yet negative sentiment in news coverage can still lead to a decline in stock prices. Conversely, optimistic sentiment can drive price increases even in the absence of strong fundamentals. Therefore, incorporating sentiment-based insights can enhance the understanding of market behavior and improve decision-making processes.

Despite its potential, financial sentiment analysis presents several domain-specific challenges. Financial text is characterized by specialized vocabulary, complex sentence structures, and context-dependent meanings. Words that are typically associated with negative sentiment in general language may carry neutral or even positive connotations in financial contexts. For example, terms such as "liability" or "debt" are not inherently negative within financial reporting. Additionally, financial narratives often include nuanced expressions, forward-looking statements, and implicit signals that require deeper contextual understanding. These complexities limit the effectiveness of generic sentiment analysis models and necessitate tailored approaches designed specifically for financial data.

To address these challenges, this research proposes a comprehensive sentiment analysis framework for financial textual data. The approach begins with the systematic collection of

relevant financial text from reliable sources, ensuring that the dataset reflects diverse market conditions and information types. The collected data undergoes rigorous preprocessing, including tokenization, stop-word removal, normalization, and, where necessary, lemmatization or stemming to standardize textual inputs. Feature extraction techniques such as bag-of-words and term frequency–inverse document frequency (TF-IDF) are employed to convert textual data into numerical representations suitable for machine learning models.

A hybrid sentiment classification strategy is adopted to improve accuracy and robustness. Lexicon-based methods are utilized to provide an interpretable baseline by assigning sentiment scores based on predefined dictionaries of financial terms. These are complemented by supervised machine learning algorithms, which learn patterns from labeled data and adapt to domain-specific language usage. The combination of these approaches aims to balance interpretability with predictive performance, overcoming the limitations associated with relying on a single method.

The system is evaluated using appropriate performance metrics, including accuracy, precision, recall, and F1-score, to assess its classification effectiveness. Experimental results demonstrate the model's capability to capture sentiment trends and distinguish between different sentiment classes with a satisfactory level of reliability. The findings highlight the importance of domain adaptation and the integration of multiple techniques in achieving meaningful results in financial text analysis.

The significance of this study lies in its contribution to the growing field of financial text mining and its practical implications for market participants. By converting unstructured textual information into actionable sentiment indicators, the proposed framework can support investors, analysts, and financial institutions in making more informed decisions. It also opens avenues for further integration with predictive models, risk assessment systems, and real-time market monitoring tools.

The remainder of this paper is structured as follows: Section II reviews existing literature related to sentiment analysis and financial text mining, highlighting key methodologies and limitations. Section III details the proposed methodology, including data collection, preprocessing, feature extraction, and model development. Section IV presents the experimental setup, results, and discussion. Finally, Section V concludes the study and outlines potential directions for future research, including the application of deep learning models and real-time sentiment analysis systems.

1.1 Background & Literature Survey

The analysis of financial markets has traditionally been grounded in structured, numerical data such as stock prices, financial ratios, and macroeconomic indicators. However, with the rapid digitization of information and the proliferation of online content, unstructured textual data has emerged as a critical source of market-relevant information. Financial news articles, corporate disclosures, earnings call transcripts, analyst reports, and even social media discussions continuously shape investor perception and influence market behavior. As a result, there has been a growing need for computational techniques capable of extracting meaningful insights from such data. Sentiment analysis, a specialized domain within Natural Language Processing (NLP), addresses this need by enabling the identification and quantification of subjective information embedded within text.

A. Foundations of Sentiment Analysis

Sentiment analysis, also referred to as opinion mining, involves the use of computational methods to determine the polarity of text—typically categorized as positive, negative, or neutral. Early developments in this field were primarily driven by applications in consumer analytics, where businesses sought to analyze customer opinions from product reviews and social media platforms. These initial approaches laid the groundwork for subsequent adaptations in more complex domains such as finance.

At its core, sentiment analysis involves several key steps: text preprocessing, feature extraction, and sentiment classification. Preprocessing typically includes tokenization, stop-word removal, normalization, and stemming or lemmatization. Feature extraction methods, such as bag-of-words and term frequency–inverse document frequency (TF-IDF), convert textual data into numerical representations. These representations are then used as input for classification algorithms that determine the sentiment polarity of the text.

While these general methodologies provide a strong foundation, their direct application to financial text is not straightforward. Financial language exhibits unique characteristics, including specialized vocabulary, context-dependent meanings, and a high degree of implicit information. Consequently, domain adaptation is essential for achieving meaningful results in financial sentiment analysis.

B. Lexicon-Based Approaches in Financial Sentiment Analysis

Lexicon-based methods represent one of the earliest approaches to sentiment analysis. These techniques rely on predefined dictionaries of words, where each word is assigned a sentiment score. The overall sentiment of a document is computed by aggregating the scores of individual words or phrases. The primary advantage of lexicon-based methods lies in their simplicity and interpretability, as the sentiment classification process is transparent and easily explainable.

In the financial domain, researchers recognized that general-purpose sentiment lexicons were insufficient due to domain-specific language usage. Words that carry negative connotations in everyday language may be neutral or even positive in financial contexts. For example, terms such as “liability,” “depreciation,” or “risk” are not inherently negative within financial reporting. To address this issue, domain-specific lexicons were developed, incorporating financial terminology and context-aware sentiment scores.

Empirical studies have demonstrated that financial lexicons significantly improve classification accuracy compared to generic dictionaries. However, lexicon-based methods are inherently limited in their ability to capture contextual nuances, syntactic relationships, and semantic dependencies. They often fail to account for negation, sarcasm, and complex sentence structures, which are common in financial narratives. As a result, while lexicon-based approaches provide a useful baseline, they are generally insufficient for achieving high-performance sentiment classification on their own.

C. Machine Learning Approaches

The limitations of lexicon-based methods led to the adoption of machine learning techniques for sentiment analysis. Supervised learning algorithms such as Naïve Bayes, Support Vector Machines (SVM), Decision Trees, and Logistic Regression became widely used due to their ability to learn patterns from labeled data. These models rely on feature extraction techniques to represent text numerically, enabling them to identify relationships between textual features and sentiment labels.

Machine learning approaches offer several advantages over lexicon-based methods. They can capture complex patterns in data, adapt to domain-specific language, and improve performance with larger datasets. In the context of financial sentiment analysis, studies have shown that machine learning models outperform lexicon-based approaches in terms of accuracy and robustness. For example, SVM models have been particularly effective due to their ability to handle high-dimensional feature spaces and sparse data.

Despite these advantages, machine learning approaches also present challenges. Their performance is highly dependent on the availability and quality of labeled datasets, which can be costly and time-consuming to obtain. Additionally, traditional machine learning models often require extensive feature engineering, which may not fully capture the semantic richness of financial text. Overfitting and model interpretability are also important considerations, particularly in financial applications where transparency is critical.

D. Deep Learning and Advanced NLP Techniques

Recent advancements in deep learning have significantly transformed the field of sentiment analysis. Neural network architectures such as Recurrent Neural Networks (RNNs), Long Short-Term Memory (LSTM) networks, and Convolutional Neural Networks (CNNs) have demonstrated strong performance in capturing sequential and contextual information in text. These models are capable of learning hierarchical representations of language, reducing the need for manual feature engineering.

In financial sentiment analysis, deep learning models have been applied to various tasks, including sentiment classification of news articles, prediction of stock price movements, and analysis of earnings call transcripts. LSTM networks, in particular, have proven effective in modeling temporal dependencies and long-range context within financial text. More recently, Transformer-based models such as BERT (Bidirectional Encoder Representations from Transformers) have achieved state-of-the-art performance in many NLP tasks, including sentiment analysis. These models leverage attention mechanisms to capture complex relationships between words, enabling a deeper understanding of textual context.

However, the application of deep learning in financial sentiment analysis is not without challenges. These models require large volumes of labeled data and significant computational resources for training. Additionally, their complexity can make them less interpretable compared to traditional methods, which may limit their adoption in certain financial applications where explainability is essential.

E. Sentiment Analysis in Financial Markets

A substantial body of literature has explored the relationship between textual sentiment and financial market behavior. Studies have shown that sentiment derived from news articles, corporate disclosures, and social media can provide valuable predictive signals for stock price movements, volatility, and trading volume. For instance, positive sentiment in earnings announcements has been associated with abnormal returns, while negative sentiment often correlates with increased market volatility.

The integration of sentiment analysis into financial models has been shown to enhance predictive performance. By combining sentiment indicators with traditional financial metrics, researchers have developed hybrid models that capture both quantitative and qualitative aspects of market behavior. These models have been applied in various contexts, including algorithmic trading, risk assessment, and portfolio management.

Social media platforms, in particular, have emerged as a rich source of real-time sentiment data. Platforms such as Twitter and financial forums provide insights into investor opinions and market sentiment at a granular level. However, the use of social media data introduces additional challenges, including noise, misinformation, and the need for robust filtering mechanisms.

F. Hybrid Approaches and Emerging Trends

To overcome the limitations of individual methods, recent research has increasingly focused on hybrid approaches that combine lexicon-based techniques with machine learning or deep learning models. These approaches aim to leverage the strengths of each method, achieving a balance between interpretability and predictive performance. For example, lexicon-based features can be used as inputs to machine learning models, enhancing their ability to capture domain-specific sentiment signals.

Another emerging trend is the use of real-time sentiment analysis systems, which process streaming data to provide up-to-date insights into market sentiment. These systems are

particularly relevant in high-frequency trading environments, where timely information is critical. Additionally, there is growing interest in multimodal analysis, where textual sentiment is combined with other data sources such as images, audio, and numerical data to provide a more comprehensive view of market behavior.

G. Research Gaps and Motivation

Despite significant advancements, several challenges remain in financial sentiment analysis. Domain specificity, data scarcity, and the dynamic nature of financial markets continue to pose difficulties for model development and evaluation. Furthermore, achieving a balance between model accuracy and interpretability remains an ongoing concern.

This study is motivated by the need to address these challenges through the development of a hybrid sentiment analysis framework tailored to financial textual data. By integrating lexicon-based methods with machine learning techniques, the proposed approach aims to achieve robust classification performance while maintaining interpretability. The research also seeks to contribute to the broader field of financial text mining by providing a structured methodology for extracting and analyzing sentiment from unstructured data.

H. Summary

In summary, the literature reveals a clear progression in sentiment analysis techniques, from simple lexicon-based methods to advanced machine learning and deep learning approaches. While each method offers distinct advantages, no single approach fully addresses the complexities of financial text. Hybrid methodologies and domain-specific adaptations have emerged as promising solutions, highlighting the importance of continued research in this area. This study builds upon these developments, aiming to provide an effective and scalable sentiment analysis framework for financial applications.

1.2 Research gap

Although sentiment analysis has matured significantly as a research domain, its application to financial textual data still exhibits several unresolved challenges that limit both theoretical advancement and practical deployment. Beyond the commonly identified issues, deeper examination of the literature reveals structural, methodological, and operational gaps that justify the need for more robust and domain-aware solutions.

A critical gap exists in the **temporal modeling of sentiment dynamics**. Most existing studies treat sentiment as a static attribute, assigning a polarity score to a document without considering how sentiment evolves over time. In financial markets, however, the timing, persistence, and transition of sentiment are crucial. For instance, a sudden shift from positive to negative sentiment may have a stronger market impact than consistently neutral sentiment. Current models often fail to capture such temporal dependencies, limiting their ability to reflect real market behavior.

Another underexplored area is the **granularity of sentiment analysis**. Many approaches operate at the document or sentence level, which can obscure fine-grained sentiment signals embedded within specific phrases or entities. Financial text often contains mixed sentiment—for example, a report may present strong revenue growth alongside concerns about future risks. Existing models may oversimplify such cases by assigning a single sentiment label, thereby losing valuable information. There is a need for more granular approaches, such as aspect-based or entity-level sentiment analysis, specifically tailored to financial contexts.

The **lack of standardized benchmarks and evaluation frameworks** also represents a significant research gap. Different studies use varying datasets, preprocessing techniques, and evaluation metrics, making it difficult to compare results or establish best practices. Moreover, many evaluations are conducted in controlled experimental settings that do not accurately reflect real-world conditions. This lack of standardization hinders the reproducibility and scalability of sentiment analysis models in financial applications.

Another important limitation is the **insufficient handling of noisy and heterogeneous data sources**. Financial text data originates from diverse channels, including formal reports, news media, and informal social platforms. Each source exhibits different linguistic characteristics, levels of reliability, and noise patterns. Existing models often focus on a single data source, limiting their generalizability. There is a gap in developing unified frameworks capable of effectively integrating and processing multi-source financial text data while filtering noise and misinformation.

The issue of **data imbalance and bias** is also prevalent in financial sentiment datasets. Positive and neutral sentiments often dominate, while negative sentiments may be underrepresented or contextually subtle. This imbalance can lead to biased models that perform poorly in detecting critical negative signals, which are often the most impactful in financial decision-making. Additionally, biases in data collection—such as overrepresentation of certain sectors or companies—can further distort model performance.

From a methodological perspective, there is a gap in the **integration of linguistic features with domain knowledge**. Many models rely heavily on statistical or neural representations of text, with limited incorporation of financial theories or expert knowledge. For example, macroeconomic indicators, industry-specific trends, and corporate governance factors are rarely integrated into sentiment analysis frameworks. Bridging this gap could lead to more context-aware and economically meaningful interpretations of sentiment.

Another emerging concern is the **interpretability of complex models**, particularly deep learning architectures. While models such as Transformers achieve high accuracy, they often function as "black boxes," making it difficult to understand how specific predictions are derived. In financial applications, where decisions may involve significant risk and regulatory scrutiny, the lack of explainability poses a major barrier to adoption. There is a clear need for approaches that provide both high performance and transparent decision-making processes.

The **adaptability of sentiment models to changing market conditions** is also insufficiently addressed. Financial markets are dynamic, with evolving language patterns, emerging terminology, and shifting investor behavior. Models trained on historical data may become outdated if they cannot adapt to new contexts. Current research offers limited solutions for continuous learning or model updating mechanisms that maintain relevance over time.

Furthermore, **real-time deployment and scalability** remain practical challenges. While many studies demonstrate promising results in offline settings, fewer address the computational and architectural requirements for deploying sentiment analysis systems in real-time environments. Issues such as latency, data throughput, and system integration are often overlooked, despite their importance in applications like algorithmic trading and live market monitoring.

Another gap lies in the **limited exploration of cross-market and cross-lingual sentiment analysis**. Most research focuses on English-language datasets and specific markets, which restricts the applicability of findings to global financial systems. In multilingual and emerging markets, such as those with diverse linguistic and economic contexts, sentiment analysis models must account for language variations and localized financial expressions.

Finally, there is a notable lack of **end-to-end frameworks that connect sentiment extraction to actionable financial outcomes**. While many studies successfully classify sentiment, fewer demonstrate how these outputs can be directly integrated into decision-making systems or predictive models. Bridging this gap requires not only accurate sentiment classification but also effective mapping of sentiment signals to financial metrics such as price movements, volatility, or risk indicators.

1.3. Research problem

Despite the growing adoption of sentiment analysis in financial domains, there remains a fundamental challenge in effectively extracting, interpreting, and utilizing sentiment from unstructured financial textual data in a manner that is both accurate and practically applicable. Financial markets are highly sensitive to information, yet a significant portion of this information

exists in textual form, making it difficult to systematically incorporate into analytical and decision-making frameworks.

The core research problem addressed in this study is the **lack of a robust, domain-adapted sentiment analysis framework capable of accurately capturing the nuanced, context-dependent nature of financial text while maintaining interpretability and practical usability**. Existing approaches exhibit notable limitations that prevent them from fully addressing this problem.

Firstly, lexicon-based methods, while interpretable, fail to capture contextual meaning and often misclassify financial terminology due to their reliance on static dictionaries. Secondly, machine learning models improve classification performance but depend heavily on labeled datasets and require extensive feature engineering, which may not generalize well across different financial contexts. Thirdly, deep learning approaches offer superior contextual understanding but introduce challenges related to computational complexity, data requirements, and lack of transparency.

In addition, financial textual data presents unique characteristics that further complicate sentiment analysis. These include domain-specific vocabulary, implicit meanings, mixed sentiment within single documents, and temporal variations in sentiment patterns. Existing models often treat sentiment as static and fail to account for its dynamic evolution over time. Moreover, the diversity of data sources—ranging from formal reports to informal social media—introduces variability and noise that are not adequately handled by conventional methods.

Another dimension of the research problem is the **difficulty in translating extracted sentiment into meaningful and actionable insights**. While many studies focus on improving sentiment classification accuracy, fewer address how these sentiment outputs can be effectively utilized in real-world financial applications. This creates a gap between theoretical model performance and practical utility.

Therefore, the problem can be formally stated as follows:
How can a sentiment analysis system be designed to accurately and efficiently extract, interpret, and quantify sentiment from financial textual data, while addressing domain-specific challenges, ensuring interpretability, and enabling meaningful application in financial decision-making?

To address this problem, this research aims to develop a hybrid sentiment analysis framework that integrates lexicon-based and machine learning techniques. The proposed solution seeks to balance accuracy and interpretability, adapt to domain-specific language, and provide a structured approach for converting unstructured financial text into actionable sentiment indicators.

Ultimately, solving this research problem contributes to bridging the gap between unstructured textual information and structured financial analysis, enabling more comprehensive and informed decision-making in modern financial environments.

1.4. Research Objectives

1.4.1 Main Objective

The main objective of this research is to **design and develop a robust, domain-adapted sentiment analysis framework capable of accurately extracting, classifying, and quantifying sentiment from financial textual data in order to generate meaningful and actionable insights**.

This objective focuses on transforming unstructured financial text—such as news articles, reports, and disclosures—into structured sentiment indicators that reflect market perception and investor behavior. The framework is intended to address key challenges associated with financial sentiment analysis, including domain-specific language, contextual ambiguity, and variability in data sources.

Furthermore, the study aims to achieve a balance between **classification accuracy and model interpretability** by integrating multiple analytical approaches within a unified system. The objective also emphasizes the development of a scalable and adaptable solution that can remain effective across different financial contexts and evolving market conditions.

Ultimately, this research seeks to establish sentiment analysis as a reliable and practical tool for financial text mining, enabling improved understanding of qualitative information and supporting more informed decision-making processes in financial environments.

1.4.2 Specific Objectives

To achieve the main objective of this research, the following specific objectives are defined:

1. **To collect and construct a relevant financial textual dataset**
   Gather financial text data from credible sources such as news articles, corporate reports, and financial disclosures, ensuring adequate coverage of different market conditions and information types.

2. **To preprocess and standardize financial textual data**
   Apply Natural Language Processing (NLP) techniques including tokenization, stop-word removal, normalization, and text cleaning to prepare the data for analysis while preserving domain-specific meaning.

3. **To develop effective feature extraction mechanisms**
   Transform textual data into numerical representations using techniques such as Bag-of-Words and Term Frequency–Inverse Document Frequency (TF-IDF) to enable machine learning-based sentiment classification.

4. **To design and implement a hybrid sentiment classification model**
   Combine lexicon-based approaches with supervised machine learning algorithms to

classify financial text into sentiment categories (positive, negative, neutral), leveraging both interpretability and predictive performance.

5. **To evaluate the performance of the sentiment analysis model**
   Assess the effectiveness of the proposed framework using standard evaluation metrics such as accuracy, precision, recall, and F1-score, ensuring robustness and reliability.

6. **To analyze sentiment patterns and trends within financial text**
   Examine how sentiment varies across different data sources and time periods, identifying meaningful patterns that reflect market perception.

7. **To ensure model adaptability and scalability**
   Develop a framework that can handle diverse financial text sources and adapt to evolving language patterns and market dynamics.

8. **To provide a structured approach for transforming unstructured data into actionable insights**
   Demonstrate how extracted sentiment can be systematically interpreted and utilized to support financial analysis and decision-making processes.

These objectives collectively guide the development of a comprehensive sentiment analysis system tailored to financial textual data, ensuring both theoretical rigor and practical applicability.

2. METHODOLOGY

2.1 Understanding the key pillars of the research domain

The foundation of this research is built upon several core pillars that collectively define the domain of financial sentiment analysis. A clear understanding of these pillars is essential to ensure that the proposed methodology is both technically sound and aligned with the complexities of financial text mining. These pillars integrate concepts from Natural Language Processing (NLP), machine learning, and financial analytics, forming the basis for the overall system design.

The first key pillar is **Natural Language Processing (NLP)**, which provides the fundamental techniques required to process and analyze unstructured textual data. Financial text is inherently complex, containing domain-specific terminology, structured narratives, and contextual dependencies. NLP enables the transformation of raw text into a structured format through processes such as tokenization, normalization, and feature extraction. This pillar is critical as it directly influences the quality of input data used in sentiment classification.

The second pillar is **Sentiment Analysis**, which focuses on identifying and quantifying the emotional tone or polarity expressed within financial text. Unlike general sentiment analysis, financial sentiment requires careful interpretation due to the presence of neutral terminology, implicit meanings, and mixed sentiment expressions. This pillar defines how textual information is converted into measurable sentiment indicators, forming the core output of the system.

The third pillar is **Machine Learning and Classification Techniques**, which are used to model and predict sentiment based on extracted features. Supervised learning algorithms enable the system to learn patterns from labeled data and generalize to unseen text. The integration of machine learning enhances the system's ability to handle complex language structures and improves classification accuracy compared to rule-based approaches.

Another important pillar is **Domain Knowledge in Finance**, which ensures that the analysis is context-aware and relevant to financial applications. Financial text often includes specialized vocabulary, industry-specific concepts, and regulatory language that require domain understanding. Incorporating financial knowledge helps in refining preprocessing steps, improving sentiment interpretation, and reducing misclassification.

The fifth pillar is **Data Preprocessing and Feature Engineering**, which plays a crucial role in preparing textual data for analysis. High-quality preprocessing ensures that noise, inconsistencies, and irrelevant information are removed, while feature engineering determines how effectively textual information is represented for model training. This pillar directly impacts the performance and reliability of the sentiment analysis system.

Additionally, **Model Evaluation and Validation** forms a key pillar, ensuring that the developed system performs reliably and generalizes well to new data. Evaluation metrics such as accuracy, precision, recall, and F1-score are used to assess classification performance, while validation techniques help prevent overfitting and ensure robustness.

Finally, the pillar of **Practical Applicability and Scalability** ensures that the developed framework is not limited to theoretical analysis but can be deployed in real-world financial environments. This includes considerations such as handling large volumes of data, adapting to evolving market conditions, and maintaining computational efficiency.

In summary, the research domain is supported by interconnected pillars: NLP, sentiment analysis, machine learning, financial domain knowledge, data preprocessing, model evaluation, and system scalability. Understanding these pillars provides a structured foundation for designing an effective sentiment analysis framework tailored to financial textual data and guides the subsequent stages of the methodology.

2.2 Financial Text Sentiment Extraction and Classification Using a Hybrid Lexicon-Based and Machine Learning Framework

The data collection phase constitutes a critical component of the proposed sentiment analysis framework, as the quality, diversity, and relevance of the collected textual data directly influence the effectiveness of the overall system. In this study, financial textual data is systematically gathered from multiple credible and publicly available sources to ensure comprehensive coverage of market-related information.

The primary data sources include **financial news articles, corporate disclosures, earnings reports, and market analysis publications**. These sources are selected due to their significant influence on investor perception and their role in shaping market sentiment. News articles provide real-time insights into market events and external factors, while corporate disclosures

and financial reports offer structured and authoritative information regarding company performance and strategic direction.

Data collection is performed using automated techniques such as web scraping and API-based retrieval, where applicable. These methods enable efficient extraction of large volumes of textual data while maintaining consistency and scalability. During this process, relevant metadata such as publication date, source, and company identifiers are also captured to facilitate temporal analysis and data organization.

To ensure data quality, several filtering criteria are applied. Irrelevant content, duplicate entries, and incomplete records are removed to minimize noise and redundancy. Additionally, only English-language text is considered to maintain consistency in subsequent Natural Language Processing (NLP) tasks. The collected data is then stored in a structured format, enabling efficient access and processing in later stages of the methodology.

Another important aspect of this phase is the preparation of labeled data for supervised learning. Where labeled datasets are not readily available, manual or semi-automated annotation techniques are employed to assign sentiment labels (positive, negative, or neutral) to selected samples. This labeled subset is essential for training and evaluating machine learning models.

Overall, the data collection process is designed to ensure that the dataset is **representative, diverse, and high-quality**, capturing a wide range of financial sentiments across different contexts and time periods. This forms a strong foundation for the subsequent preprocessing, feature extraction, and sentiment classification stages of the research.

2.2.2 Methodology

This study adopts a **hybrid, multi-stage sentiment analysis framework** specifically designed for financial textual data. The methodology is inspired by advanced financial NLP architectures and incorporates both machine learning and lexicon-based techniques to address domain-specific challenges such as contextual ambiguity, class imbalance, and interpretability.

A. Problem Formulation and Labeling Strategy

Unlike traditional sentiment classification, which directly assigns polarity labels, this study reformulates sentiment detection as a **two-stage learning problem** to improve robustness and directional accuracy.

In financial contexts, not all textual information significantly impacts market behavior. Therefore, the first stage focuses on identifying whether a given text is **market-relevant (impactful)**, while the second stage determines the **direction of sentiment (positive, negative, neutral)**.

This approach is motivated by findings in prior architectures where single-stage classification resulted in poor performance due to class imbalance and dominance of neutral samples

B. Text Representation and Embedding

To capture semantic meaning beyond surface-level text, financial documents are transformed into **vector representations using embedding techniques**. These embeddings encode contextual relationships between words and enable the model to understand financial narratives more effectively.

The process includes:

- Converting preprocessed text into dense vector representations
- Capturing semantic similarity between financial events
- Enabling downstream classification tasks

This representation allows the model to move beyond simple keyword matching and instead understand **contextual sentiment**, which is critical in financial text.

C. Stage 1: Market Impact Detection Model

The first stage of the model is designed to determine whether a piece of financial text has a **significant impact on market behavior**.

A supervised machine learning model (e.g., gradient boosting such as XGBoost) is trained to estimate:

- Whether a news item or document is **market-moving** or not

This stage acts as a **filtering mechanism**, ensuring that only relevant financial text proceeds to sentiment classification.

To improve sensitivity:

- A **recall-oriented threshold** is applied
- The model prioritizes detecting impactful events, even at the cost of some false positives

This design choice is crucial because missing impactful information is more costly than misclassifying non-impactful text.

D. Stage 2: Sentiment Direction Classification

For text identified as impactful, the second stage performs **sentiment classification**, predicting whether the sentiment is:

- Positive
- Negative
- Neutral

This stage uses supervised learning models trained on labeled financial text data. However, as identified in prior work, directional prediction is inherently more difficult than impact detection, with models often showing near-random performance when used alone .

To address this limitation, additional enhancements are introduced.

E. Lexicon-Based Sentiment Enhancement

To improve interpretability and robustness, a **domain-specific financial lexicon** is integrated into the model.

The lexicon assigns sentiment scores to financial terms, and a document-level score is computed as:

- Aggregation of individual word sentiment values
- Adjustment for financial context and polarity

This component is particularly useful when:

- Model confidence is low
- Text contains explicit sentiment indicators

As highlighted in the reference architecture, lexicon-based scoring is used as a **fallback or complementary mechanism** when model predictions are uncertain

F. Hybrid Decision Mechanism

A key innovation in the methodology is the **adaptive combination of machine learning predictions and lexicon-based scores**.

Instead of relying solely on one method, the final sentiment prediction is obtained through:

- Model probability outputs
- Lexicon-derived sentiment confidence
- A weighted fusion mechanism

This hybrid approach:

- Improves classification accuracy
- Maintains interpretability
- Reduces dependency on large labeled datasets

G. Handling Class Imbalance

Financial sentiment datasets typically exhibit **class imbalance**, where neutral sentiment dominates.

To mitigate this issue:

- The problem is decomposed into two stages (impact + direction)
- Threshold tuning is applied
- Model evaluation focuses on **F1-score and recall**, rather than accuracy alone

This restructuring significantly improves the model's ability to detect minority classes (positive and negative sentiment).

H. Semantic Similarity and Context Retrieval (Optional Enhancement)

To enhance interpretability and contextual understanding, the methodology incorporates a **semantic similarity mechanism**:

- Text embeddings are compared using similarity measures (e.g., cosine similarity)
- Historical documents with similar semantic meaning are retrieved

As described in the reference system, this enables:

- Identification of past events with similar sentiment patterns
- Contextual explanation of model predictions

This component adds an explainability layer by linking predictions to real historical examples.

I. Model Evaluation

The performance of the sentiment model is evaluated using multiple metrics:

- **Accuracy** – overall correctness
- **Precision** – correctness of positive predictions
- **Recall** – ability to detect relevant sentiment
- **F1-score** – balance between precision and recall
- **ROC-AUC** – classification capability across thresholds

Special emphasis is placed on:

- Performance on minority classes
- Robustness of directional predictions

J. Summary of the Sentiment Model Pipeline

The complete sentiment analysis workflow is summarized as:

Text Input → Preprocessing → Embedding → Stage 1 (Impact Detection) → Stage 2 (Sentiment Classification) → Lexicon Enhancement → Hybrid Decision → Output Sentiment

---

K. Key Contributions of the Methodology

The proposed sentiment model introduces several improvements over traditional approaches:

- A **two-stage architecture** to handle imbalance and improve robustness
- A **hybrid model** combining machine learning and lexicon methods
- Enhanced **interpretability through lexicon scoring and similarity retrieval**
- Adaptation to **financial domain-specific language and context**

2.4 High-Level Architecture Diagram

2.5 Project Requirements that have been achieved

2.5.1 Functional Requirements

1. Financial Text Data Ingestion
The system is capable of collecting and ingesting financial textual data from multiple sources such as news articles, reports, and disclosures. It supports automated data retrieval and ensures that relevant textual information is available for analysis.

---

2. Text Preprocessing and Cleaning
The system performs comprehensive preprocessing of raw textual data, including noise removal, tokenization, normalization, and stop-word elimination. This ensures that the input data is standardized and suitable for further processing.

---

3. Feature Extraction and Representation
The system converts processed text into numerical representations using techniques such as TF-IDF and n-grams. This enables effective input for machine learning models and improves classification performance.

---

4. Market Impact Detection
The system implements a first-stage classification model to identify whether a given piece of financial text has a significant impact on market behavior. This filtering mechanism ensures that only relevant information is considered for sentiment analysis.

---

5. Sentiment Classification
The system classifies financial text into sentiment categories (positive, negative, and neutral)

using supervised machine learning models. This provides structured sentiment outputs from unstructured text data.

---

6. Lexicon-Based Sentiment Scoring
A domain-specific financial lexicon is integrated into the system to compute sentiment scores based on predefined word polarity. This enhances interpretability and supports model predictions.

---

7. Hybrid Sentiment Decision Mechanism
The system combines machine learning predictions with lexicon-based scores to produce final sentiment outputs. This hybrid approach improves robustness and accuracy.

---

8. Handling Class Imbalance
The system incorporates a two-stage classification approach and threshold tuning to address class imbalance issues commonly found in financial datasets.

---

9. Model Evaluation and Validation
The system evaluates model performance using metrics such as accuracy, precision, recall, F1-score, and ROC-AUC. This ensures reliability and effectiveness of the sentiment classification process.

---

10. Sentiment Output Generation
The system produces structured sentiment outputs that can be used for analysis, visualization, or integration into decision-making processes.

2.5.2 Non-Functional Requirements

In addition to functional capabilities, the developed sentiment analysis system satisfies several non-functional requirements that ensure the system operates efficiently, reliably, and is suitable for practical deployment. These requirements define the quality attributes of the system and support its overall performance and usability.

---

1. Performance Efficiency
The system is designed to process large volumes of financial textual data within a reasonable

time frame. Efficient preprocessing pipelines and optimized machine learning models ensure low latency in sentiment prediction, making the system suitable for near real-time analysis.

---

2. Scalability
The architecture supports scalability to handle increasing data volumes and additional data sources. The modular design allows components such as data ingestion, preprocessing, and classification to be scaled independently based on workload requirements.

---

3. Accuracy and Reliability
The system maintains high classification accuracy through the use of a hybrid model combining machine learning and lexicon-based approaches. Robust evaluation techniques and validation methods ensure consistent and reliable performance across different datasets.

---

4. Interpretability and Explainability
Given the importance of transparency in financial applications, the system provides interpretable outputs. Lexicon-based scoring and model insights enable users to understand how sentiment decisions are derived, supporting trust and accountability.

---

5. Maintainability
The system is structured using a modular design, allowing easy updates, debugging, and extension of individual components. This ensures that improvements or modifications can be implemented without affecting the entire system.

---

6. Flexibility and Adaptability
The system is adaptable to different types of financial text sources and evolving language patterns. New data sources, updated lexicons, or improved models can be integrated with minimal restructuring.

---

7. Robustness
The system is capable of handling noisy, incomplete, or inconsistent textual data. Preprocessing techniques and validation mechanisms ensure stable performance even under imperfect data conditions.

---

8. Usability
The system produces clear and structured outputs that are easy to interpret and utilize. This ensures that users, including analysts and researchers, can effectively leverage sentiment insights without requiring deep technical expertise.

---

9. Resource Efficiency
The methodology balances computational complexity and performance by combining lightweight lexicon methods with more advanced machine learning models. This ensures efficient use of computational resources.

---

10. Reproducibility
The system ensures that results can be consistently reproduced through well-defined preprocessing steps, model configurations, and evaluation procedures. This is essential for research validity and future improvements.

---

11. Security and Data Integrity
Collected data is handled in a controlled manner, ensuring that data integrity is maintained throughout the processing pipeline. Proper data handling mechanisms prevent corruption or loss of information.

---

12. Extensibility
The system is designed to support future enhancements such as integration with deep learning models, real-time data streams, or additional financial analytics modules.

In summary, the non-functional requirements achieved by the system ensure that it is not only technically effective but also practical, scalable, and reliable for real-world financial sentiment analysis applications.

2.5.3 Other Requirements

In addition to the functional and non-functional requirements, several supporting requirements were identified and achieved to ensure the successful development, deployment, and usability of the sentiment analysis system within a research and practical context.

---

1. Data Availability and Accessibility
The system requires access to reliable and up-to-date financial textual data sources. Publicly available datasets and financial news platforms are utilized to ensure continuous data availability. The system is designed to accommodate different data formats and sources without significant restructuring.

---

2. Development Environment Requirements
The implementation is carried out using standard development tools and programming environments suitable for machine learning and NLP tasks. This includes support for libraries and frameworks required for text processing, model training, and evaluation.

---

3. Hardware Requirements
The system is designed to operate on standard computing hardware without requiring specialized high-performance infrastructure. However, for larger datasets or advanced models, higher computational resources (e.g., increased RAM or GPU support) may enhance performance.

---

4. Software Dependencies
The system depends on widely used software libraries for Natural Language Processing and machine learning. These dependencies are selected based on stability, community support, and compatibility to ensure long-term usability and maintainability.

---

5. Data Storage and Management
Efficient storage mechanisms are implemented to manage collected textual data and processed outputs. Structured storage formats are used to enable easy retrieval, updating, and analysis of data.

---

6. Integration Capability
The system is designed to allow integration with other analytical modules or external systems. This enables the sentiment outputs to be used in broader applications such as financial forecasting, dashboards, or decision-support systems.

---

7. Documentation and Reproducibility
Comprehensive documentation is maintained for all stages of the system, including data

collection, preprocessing, model development, and evaluation. This ensures that the research can be reproduced and extended in future work.

---

8. Ethical Considerations
The system uses publicly available data and adheres to ethical standards in data usage. Care is taken to avoid misuse of information and ensure that the analysis does not violate any data privacy or usage policies.

---

9. Testing and Validation Requirements
The system undergoes systematic testing to verify correctness, stability, and performance. Both unit-level and system-level validations are conducted to ensure that each component functions as expected.

---

10. Deployment Considerations
Although primarily developed for research purposes, the system is designed with potential deployment in mind. This includes considerations for scalability, modular integration, and adaptability to real-world financial environments.

---

11. Version Control and Collaboration
Version control mechanisms are utilized to manage code changes and support collaborative development. This ensures traceability, consistency, and efficient teamwork throughout the project lifecycle.

---

12. Future Enhancement Support
The system is designed to support future improvements, including the incorporation of deep learning models, real-time data processing, and advanced analytics. This ensures that the system remains relevant and extendable beyond the current scope.

2.6. Consideration of the aspects of the system

2.6.1 Social aspects

The development and deployment of a financial sentiment analysis system carry several social implications, particularly in how information is interpreted, disseminated, and utilized by individuals and institutions. These aspects are important to ensure that the system contributes positively to society while minimizing potential risks.

One key social aspect is the **democratization of financial information**. Traditionally, access to advanced analytical tools and market insights has been limited to institutional investors with significant resources. The proposed sentiment analysis system helps bridge this gap by transforming complex financial text into understandable sentiment indicators. This enables retail investors and smaller stakeholders to access insights that were previously difficult to obtain, thereby promoting more inclusive participation in financial markets.

Another important consideration is the **influence on investor behavior**. Sentiment analysis systems can shape decision-making by highlighting positive or negative market signals. While this can support more informed decisions, there is also a risk that users may over-rely on automated sentiment outputs without critical evaluation. This could lead to herd behavior or amplified market reactions. Therefore, the system is designed to provide **supportive insights rather than prescriptive decisions**, encouraging users to combine sentiment analysis with other forms of financial evaluation.

The issue of **information credibility and bias** is also significant. Financial text data may contain biased, incomplete, or misleading information depending on the source. If not properly managed, these biases can propagate through the sentiment analysis model and influence outcomes. To mitigate this, the system incorporates data filtering and preprocessing mechanisms, and emphasizes the use of credible sources. However, users must remain aware that sentiment outputs reflect the input data and are not absolute truths.

Transparency and interpretability play a crucial role in building user trust. In financial contexts, users need to understand how conclusions are derived. The integration of lexicon-based scoring alongside machine learning models enhances explainability, allowing users to see the contributing factors behind sentiment classifications. This reduces the perception of the system as a “black box” and promotes responsible usage.

Another social consideration is the **impact on employment and professional roles**. The adoption of automated sentiment analysis tools may reduce reliance on manual analysis performed by financial analysts. However, rather than replacing human expertise, the system is intended to augment analytical capabilities by handling large volumes of data efficiently, allowing professionals to focus on higher-level decision-making.

The system also considers **ethical data usage**. Only publicly available financial text is utilized, and no personal or sensitive data is processed. This ensures compliance with ethical standards and avoids privacy concerns. Additionally, the system does not manipulate or generate misleading information but strictly analyzes existing data.

Finally, there is the aspect of **financial literacy and accessibility**. By presenting sentiment insights in a structured and understandable manner, the system can contribute to improving users’ understanding of market dynamics. However, care must be taken to ensure that outputs are not oversimplified in a way that misleads users about the complexity of financial markets.

In summary, the social aspects of the system highlight its potential to enhance accessibility, support informed decision-making, and improve financial literacy, while also acknowledging

risks related to bias, over-reliance, and information interpretation. Addressing these considerations ensures that the system is used responsibly and contributes positively to the broader financial ecosystem.

Security is a critical consideration in the design and implementation of the sentiment analysis system, particularly due to its reliance on data collection, processing, and storage. Although the system primarily operates on publicly available financial textual data, ensuring the integrity, confidentiality, and reliability of the system remains essential.

One of the primary security concerns is **data integrity**. The system must ensure that the collected financial text data is not altered, corrupted, or tampered with during collection, storage, or processing. To address this, validation mechanisms are applied during data ingestion, including filtering malformed inputs and verifying data consistency. This ensures that sentiment analysis results are based on accurate and trustworthy data.

Another important aspect is **secure data handling and storage**. The system organizes and stores textual data and processed outputs in a structured manner to prevent data loss or unauthorized modification. Access to stored data is controlled to ensure that only authorized processes or users can retrieve or modify information. Proper backup mechanisms are also considered to safeguard against accidental data loss.

The system also addresses **protection against malicious input and attacks**. Since data may be collected from external sources, there is a risk of encountering harmful or manipulated content. Input validation and preprocessing steps are designed to detect and filter out potentially harmful data, such as scripts or malformed text, thereby reducing the risk of system vulnerabilities such as injection attacks.

System reliability and availability are also closely tied to security. The system is designed to operate consistently without interruptions, even when processing large volumes of data. Mechanisms such as error handling and fault tolerance are implemented to ensure that failures in one component do not compromise the entire system.

Another key consideration is **model security and robustness**. Machine learning models can be vulnerable to adversarial inputs or biased data that may affect predictions. While the system primarily operates in a research context, care is taken to ensure that the model is trained on clean and representative data, reducing the risk of manipulation or misleading outputs.

Access control and authorization are also important, especially if the system is extended to a deployed environment. Role-based access mechanisms can be implemented to restrict system functionalities based on user roles, ensuring that sensitive operations such as model updates or data modification are controlled.

In addition, **secure integration and communication** are considered when the system interacts with external data sources or other systems. Ensuring safe data transfer mechanisms helps prevent unauthorized interception or data leakage.

Finally, the system adheres to **ethical and responsible data usage practices**, ensuring that only publicly available and non-sensitive financial data is used. This minimizes privacy risks and aligns the system with acceptable data governance standards.

In summary, the security aspects of the system focus on maintaining data integrity, protecting against malicious inputs, ensuring reliable operation, and enabling controlled access. These measures collectively contribute to building a secure and trustworthy sentiment analysis system suitable for financial applications.

2.6.3 Ethical aspects

The development and deployment of the financial sentiment analysis system involve several ethical considerations to ensure that the system is used responsibly, transparently, and without causing harm to users or the broader financial ecosystem. These considerations are particularly important given the potential influence of sentiment-based insights on investment decisions and market behavior.

One of the primary ethical aspects is **responsible use of data**. The system relies exclusively on publicly available financial textual data such as news articles and corporate disclosures. No personal, private, or sensitive data is collected or processed. This ensures compliance with ethical data usage standards and avoids privacy violations. Additionally, proper attribution of data sources is maintained where applicable.

Another important consideration is **bias and fairness**. Financial text data may contain inherent biases depending on the source, publication style, or market perspective. If not properly addressed, these biases can propagate through the sentiment analysis model and lead to skewed or misleading results. To mitigate this, the system incorporates preprocessing and filtering mechanisms and uses diverse data sources where possible. However, users are informed that the outputs reflect the underlying data and may not represent an objective truth.

The issue of **transparency and explainability** is also central to ethical system design. In financial applications, users must understand how conclusions are derived. The system addresses this by integrating interpretable components such as lexicon-based sentiment scoring alongside machine learning models. This allows users to trace the reasoning behind sentiment classifications, reducing reliance on opaque “black-box” decisions.

Another key ethical concern is **misuse and over-reliance**. There is a risk that users may treat sentiment outputs as definitive investment advice rather than analytical support. This could lead to poor financial decisions or contribute to herd behavior in markets. To address this, the system is designed as a **decision-support tool**, not a decision-making authority. Users are expected to combine sentiment insights with other forms of financial analysis and professional judgment.

The system also considers the **impact on market dynamics**. Large-scale use of sentiment analysis tools could amplify reactions to news and potentially contribute to market volatility. While the system itself does not directly influence markets, its outputs could indirectly affect

user behavior. Therefore, it is important that the system is used responsibly and that its limitations are clearly communicated.

Accountability is another important aspect. The system provides analytical outputs based on data and models, but it does not assume responsibility for decisions made by users based on those outputs. Clear documentation and explanation of system limitations are provided to ensure that users understand the scope and boundaries of the system.

Additionally, the system adheres to **academic integrity and research ethics**. All methodologies, data handling procedures, and results are documented transparently to ensure reproducibility and credibility. Any external sources or references are properly acknowledged.

Finally, the system promotes **ethical innovation**, aiming to enhance financial analysis without replacing human judgment or encouraging irresponsible behavior. It is designed to augment human decision-making rather than automate it entirely.

In summary, the ethical aspects of the system focus on responsible data usage, bias mitigation, transparency, prevention of misuse, and accountability. Addressing these considerations ensures that the sentiment analysis system is aligned with ethical standards and contributes positively to both research and practical financial applications.